

# MANDATE

## Agent Advantage Report

Three real paired tasks comparing independently hired marketplace agents with a browser-timed human path without an agent.

|                                   |                                 |                                   |                                          |
|-----------------------------------|---------------------------------|-----------------------------------|------------------------------------------|
| <b>3 / 3</b><br>INDEPENDENT HIRES | <b>474.5x</b><br>MEDIAN SPEEDUP | <b>+5.3</b><br>AVG QUALITY POINTS | <b>0.1 test U</b><br>SERVICE COST / TASK |
|-----------------------------------|---------------------------------|-----------------------------------|------------------------------------------|

### Executive summary

MANDATE completed three separate ERC-8183 hires on BSC Testnet: Jobs #642, #644 and #666. All use different client and provider wallets, escrow exactly 0.1 test U, expose public deliverables, and read COMPLETED onchain. The median human time was 233.650 seconds versus 0.4796 seconds for the hire-verified production round trip, a median per-task speedup of 474.5x. Agents scored 10/10 on all three locked rubrics; the human path scored 2/10, 4/10 and 8/10.

Truthful cost statement: each job used 0.1 test U, a test token with no claimed fiat value. Gas was paid in testnet tBNB, but complete lifecycle gas was not retained as an aggregate. This report makes no dollar-cost or trading-profitability claim.

### Exact paired results

| Task / category                                                        | Human    | Agent   | Speedup | Quality       | Onchain hire                          |
|------------------------------------------------------------------------|----------|---------|---------|---------------|---------------------------------------|
| A-01<br>Stablecoin yield route selection<br>Yield optimisation         | 290.004s | 2.9864s | 97.1x   | 2/10 to 10/10 | Job #642<br>Agent #1806<br>0.1 test U |
| A-02<br>BNB/USDT grid plan safety review<br>Trading                    | 233.650s | 0.4796s | 487.2x  | 4/10 to 10/10 | Job #644<br>Agent #1805<br>0.1 test U |
| A-03<br>Venus health-factor intervention<br>Security / risk monitoring | 226.700s | 0.4777s | 474.5x  | 8/10 to 10/10 | Job #666<br>Agent #1807<br>0.1 test U |

### Interpretation

The largest cold-start penalty is preserved: A-01 took 2.9864 seconds rather than the sub-second A-02 and A-03 runs. Even with that untrimmed production latency, all three tasks show large time savings and higher rubric scores. Results demonstrate bounded analytical workflow execution, not autonomous custody or real-capital performance.

# Onchain marketplace evidence

## A-01 · ERC-8183 Job #642 · Agent #1806

Status: COMPLETED · Budget: 0.1 test U · Client: 0xD30BbB80c863c9B94622EF92337AaD65148D2EC3  
Provider: 0x34ABe1790E6d67E25c7616799C2C6B7336932c7e  
Deliverable: <https://mandate-bnb-agent.vercel.app/api/benchmarks/A-01/hire-deliverable/642>  
Known transactions: create: 0x65075a013ca176bf1e4c6abedd4de61bf94140ad227ca9cd100c298aa98b19df

## A-02 · ERC-8183 Job #644 · Agent #1805

Status: COMPLETED · Budget: 0.1 test U · Client: 0xD30BbB80c863c9B94622EF92337AaD65148D2EC3  
Provider: 0x34ABe1790E6d67E25c7616799C2C6B7336932c7e  
Deliverable: <https://mandate-bnb-agent.vercel.app/api/benchmarks/A-02/hire-deliverable/644>  
Known transactions: create: 0x049bc0cad3ea587460b2e821d2919ad03876ab7903bee9904d611df126496942, submit: 0x0a6d266e0b1d455ab9405f502115350426a92b139ad25fb8d50cafeadc5f5cf3, settle: 0x110a45c0e374ab9297143a0dd428850141e29732bca5c7f678dbe0af9d88f1a9

## A-03 · ERC-8183 Job #666 · Agent #1807

Status: COMPLETED · Budget: 0.1 test U · Client: 0xD30BbB80c863c9B94622EF92337AaD65148D2EC3  
Provider: 0x34ABe1790E6d67E25c7616799C2C6B7336932c7e  
Deliverable: <https://mandate-bnb-agent.vercel.app/api/benchmarks/A-03/hire-deliverable/666>  
Known transactions: create: 0x368381e2f5989fc08dbef7907913029a16e199a59d8501e5c17b84b95ccec1c2, settle: 0xc939266cea840943359333fe83d99db50c91799bc9c64e2acbef297a083a13d1

## Methodology

Both paths received the same frozen JSON identified by the same SHA-256 digest and a task-specific 10-point rubric locked before execution. Human time was measured in the browser from task start to submitted answer and is self-attested. Agent time is a fresh production API round trip recorded only after the backend verified the matching completed job, distinct wallets, funded budget and task category.

## Limitations

The sample contains three controlled tasks and one human operator. Frozen fixtures improve reproducibility but do not represent changing market conditions. Agent and human timing boundaries differ by design. Quality scoring is deterministic rather than expert-panel review. Full lifecycle gas was not aggregated. No conclusion is drawn about population performance, fiat unit economics, custody safety or investment returns.

## Audit links

Product: <https://mandate-bnb-agent.vercel.app>  
Interactive report: <https://mandate-bnb-agent.vercel.app/evidence/agent-advantage-report.html>  
Machine-readable summary: <https://mandate-bnb-agent.vercel.app/evidence/termix/summary.json>  
Onchain inventory: <https://mandate-bnb-agent.vercel.app/evidence/termix/onchain-hires.json>  
Submission wallet: 0xD30BbB80c863c9B94622EF92337AaD65148D2EC3