

MANDATE

Agent Advantage Report

Three real, same-input tasks comparing marketplace agents with a browser-timed human path without an agent.

3 / 3 PAIRED TASKS	866.7x MEDIAN SPEEDUP	+5.3 pts AVG QUALITY LIFT	\$0.00 OBSERVED COST GAP
------------------------------	---------------------------------	-------------------------------------	------------------------------------

Executive summary

MANDATE completed all three paired tasks required by the TermiX track, including trading and security. Each pair used the same frozen input SHA-256 and a rubric locked before either path ran. The median human time was 233.650 seconds; the median agent browser round trip was 0.2718 seconds. The median per-task speedup was 866.7x. Agents scored 10/10 on all tasks, while human scores were 2/10, 4/10, and 8/10.

Observed incremental cash cost was \$0 for both paths. This report therefore claims a measured speed and quality advantage, not a paid-service cost advantage.

Quality on the same locked rubric




Figure 1. Quality scores on the pre-committed 10-point rubric. No results were substituted or manually rescored after either path ran.

Exact paired results

Human time is browser timed from task start to submitted answer. Agent time is the browser-observed round trip to the public production API.

Task	Category	Human	Agent	Speedup	Quality	Cash cost
A-01 Stablecoin yield route selection	Yield optimisation	290.004 s	0.2718 s	1067.0x	2/10 -> 10/10	\$0 -> \$0
A-02 BNB/USDT grid plan safety review	Trading	233.650 s	0.2696 s	866.7x	4/10 -> 10/10	\$0 -> \$0
A-03 Venus health-factor intervention	Security / risk monitoring	226.700 s	0.2908 s	779.6x	8/10 -> 10/10	\$0 -> \$0

A-01 - Stablecoin yield route selection

Category: Yield optimisation. Time saved: 289.7322 seconds. Quality lift: +8 points. Frozen input SHA-256:

2150d74ff5e31d5d420eff2435cd6dc8030845c5bc84db68fd4af02a817aee21

A-02 - BNB/USDT grid plan safety review

Category: Trading. Time saved: 233.3804 seconds. Quality lift: +6 points. Frozen input SHA-256:

153f5f9a569877b6494132642c256c45ad864e1a14ae5263fca11827ac2c7eb2

A-03 - Venus health-factor intervention

Category: Security / risk monitoring. Time saved: 226.4092 seconds. Quality lift: +2 points. Frozen input SHA-256:

a398c50ddcc3cc105b602e7052e5213eef35f8f6bbaff69a7cde11047f5bc27e

Cost interpretation

Observed incremental cash cost was \$0.00 on both paths. For sensitivity only, valuing the measured human time at a disclosed \$15/hour yields \$1.2084, \$0.9735, and \$0.9446 across A-01 to A-03 (\$3.1265 total). This is modeled opportunity cost, not money paid. Agent hosting and development are excluded, and no marketplace service fee was charged during these runs. A future paid rerun is required before claiming a monetary advantage.

Method and audit trail

Controls

Both paths received the exact JSON represented by the same input digest. The rubric was visible and locked before either path started. Raw outputs retain timestamps, scorer version, timings, decisions, rejected options, calculations, and recommendations. A-02 includes a verification note for schema-preserving numeric normalization (18.0 in Python versus 18 in JSON); the value is unchanged and the reconstructed hash matches.

Evidence inventory

Task	Human output	Agent output	Hash status
A-01	evidence/termix/raw/A-01-baseline.json	evidence/termix/raw/A-01-agent.json	Exact match
A-02	evidence/termix/raw/A-02-baseline.json	evidence/termix/raw/A-02-agent.json	Verified with numeric type normalization
A-03	evidence/termix/raw/A-03-baseline.json	evidence/termix/raw/A-03-agent.json	Exact match

Separate onchain proof

ERC-8183 YieldRoute Job #506 completed a seven-step BSC Testnet lifecycle with exact 0.1 test U escrow and zero residual allowance. Its submitted deliverable binds a public live-data report and evidence SHA-256. This proves marketplace settlement mechanics separately; it was not the payment rail for these three benchmark runs.

Limitations

The human sample is one operator and three tasks. Agent timings are API round trips, while human timings include reading and composition. Market inputs were frozen fixtures, not live capital decisions. The quality rubric is deterministic and task-specific. These controls support reproducibility but do not estimate population-wide performance or paid unit economics.

Sources

Official TermiX track: <https://www.bnbchain.org/en/hackathons/smart-money-era>

Onchain settlement: <https://testnet.bscscan.com/tx/0xf423d6403c8e7926ea0e125c3b216226b95856fc836293645ef14c8ae531f043>

Evidence Passport: <https://mandate-bnb-agent.vercel.app/evidence/evidence-passport-506.json>

Public product: <https://mandate-bnb-agent.vercel.app>

Submission wallet: 0xD30BbB80c863c9B94622EF92337AaD65148D2EC3